


The background of the slide is a light gray gradient. It is decorated with numerous realistic water droplets of various sizes. Some droplets are at the top left, some are in the center, and a larger cluster is at the bottom right. Each droplet has a highlight and a shadow, giving it a 3D appearance.

# SKIN CANCER DETECTION MODELS

CLASSIFY CANCER IMAGES AS MALIGNANT OR BENIGN

# DATA

- TRAINING DATAT SET: 2000 DERMOSCOPIC IMAGES
  - CLASS LABELS – MALIGNANT/BENIGN (1 /0) FOR MELANOMA
  - 374 MALIGNANT
  - 1626 BENIGN
- VALIDATION DATA SET: 150 DERMOSCOPIC IMAGES
  - CLASS LABELS – MALIGNANT/BENIGN (1 /0) FOR MELANOMA
  - 30 MALIGNANT
  - 120 BENIGN



# DATA PREPARATION



- RESIZING OF IMAGES - CONVERTED TO 224 X 224 X 3
- ROTATION OF IMAGES – 180 AND 90 DEGREES
- UPSCALING OF MALIGNANT IMAGES
- CONVERTED IMAGE PIXELS TO A VECTOR
- DEVELOPED TRAIN DATA (.H5 FILES) – PREDICTOR MATRIX AND LABEL VECTOR
- STANDARDIZED PREDICTORS

# MODELS DEVELOPED

- I. DEEP MODELS USING KERAS LIBRARIES
- II. SHALLOW NEURAL NETWORK – 2 LAYERS
- III. DEEP NEURAL NETWORK – 5 LAYERS

OPTIMIZER = Adam()  
METRICS=['accuracy']  
LOSS = 'binary\_crossentropy'  
KERNEL\_INITIAL = 'glorot\_uniform'

# CCN MODEL 1





Epoch 2/10  
loss: 2.9977  
acc: 0.8130

deepPredict(model1,x\_test,y\_test)  
Test score: 3.20604784648  
Test accuracy: 0.8000000000795

OPTIMIZER = Adam()  
METRICS = ['accuracy']  
LOSS = 'binary\_crossentropy'  
KERNEL\_INITIAL = 'glorot\_uniform'

## CCN MODEL 2






Epoch 2/10  
loss: 2.9977  
acc: 0.8130

deepPredict(model1,x\_test,y\_test)  
Test score: 3.20604784648  
Test accuracy: 0.8000000000795

OPTIMIZER = Adam()  
METRICS = ['accuracy']  
LOSS = 'binary\_crossentropy'  
KERNEL\_INITIAL = 'glorot\_uniform'

## CCN MODEL 3





Epoch 10/10  
loss: 5.0137  
acc: 0.6872

deepPredict(model1,x\_test,y\_test)  
Test score: 3.20604784648  
Test accuracy: 0.8000000000795



OPTIMIZER = rmsprop  
METRICS = ['accuracy']  
LOSS = 'binary\_crossentropy'

# CCN MODEL 4



Epoch 10/10

loss: 0.4868

acc: 0.8183

`deepPredict(model1,x_test,y_test)`



Test score: 0.488429104487

Test accuracy: 0.800000000795



OPTIMIZER = rmsprop  
METRICS = ['accuracy']  
LOSS = 'binary\_crossentropy'

# CCN MODEL 4 IMAGES ROTATED (180)



Epoch 13/14

loss: 0.5415

acc: 0.7401



`deepPredict(model1,x_test,y_test)`

Test score: 0.5447554938

Test accuracy: 0.7699999996026

OPTIMIZER = rmsprop  
METRICS = ['accuracy']  
LOSS = 'binary\_crossentropy'

# CCN MODEL 4 IMAGES ROTATED (90)



Epoch 13/14

loss: 0.5433

acc: 0.7377

`deepPredict(model1,x_test,y_test)`

Test score: 0.527270306746

Test accuracy: 0.77666667064

# SHALLOW NEURAL NETWORK

Cost after iteration 0: 0.6960335407116954

Cost after iteration 100: 0.4724549661840295

Cost after iteration 200: 0.4695814213451869

---

Cost after iteration 2200: 0.44743597272294433


Cost after iteration 2300: 0.4461306727321353

Cost after iteration 2400: 0.44582365392664597

`predictions_train = predict(X, Y, parameters)`

**Accuracy: 0.82**

for shallow 2 layer model [150528, 7, 1]



# DEEP NEURAL NETWORK

Cost after iteration 0: 0.653496

Cost after iteration 100: 0.473703

Cost after iteration 200: 0.465915

---

---

Cost after iteration 2200: 0.452765


Cost after iteration 2300: 0.450670

Cost after iteration 2400: 0.450000

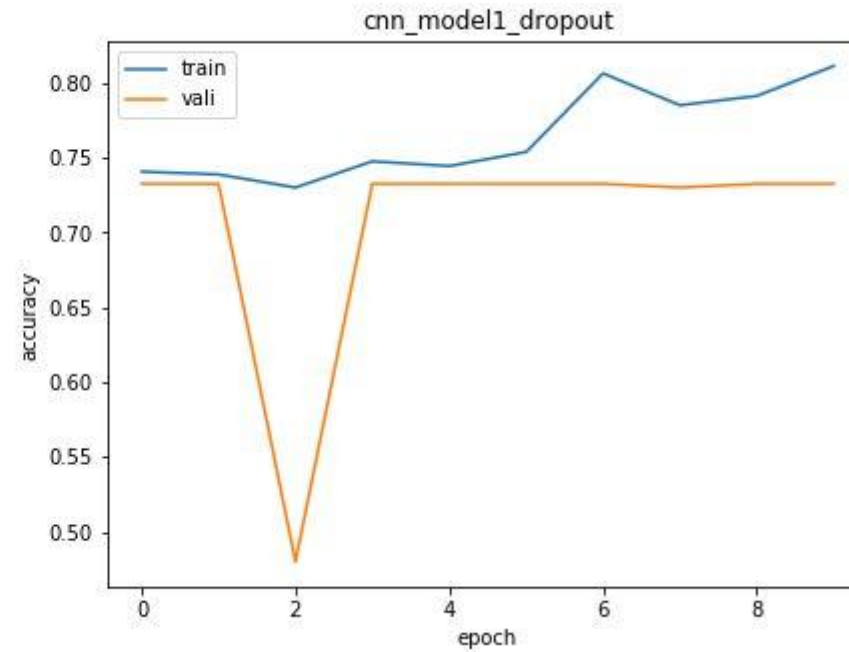

`pred_train = predict(X, Y, parameters)`

**Accuracy: 0.8145**

for [150528, 20, 7, 5, 1] 5-layer model



# KERAS NEURAL NETWORK\_1 RELU, 1 SOFTMAX



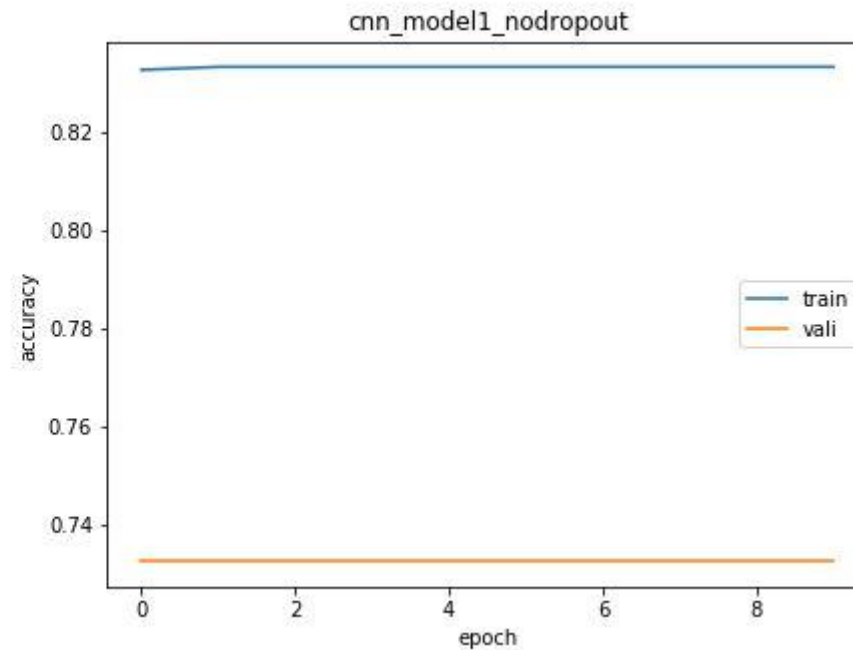

## With dropout

Epoch 4/10

```
1550/1600 [=====>.] - ETA: 3s - loss: 1.2906 - acc: 0.7439 Epoch 00003: val_acc improved from 0.73250 to 0.73250. saving model to model1 dropout 1 03 0.73.hdf5
```



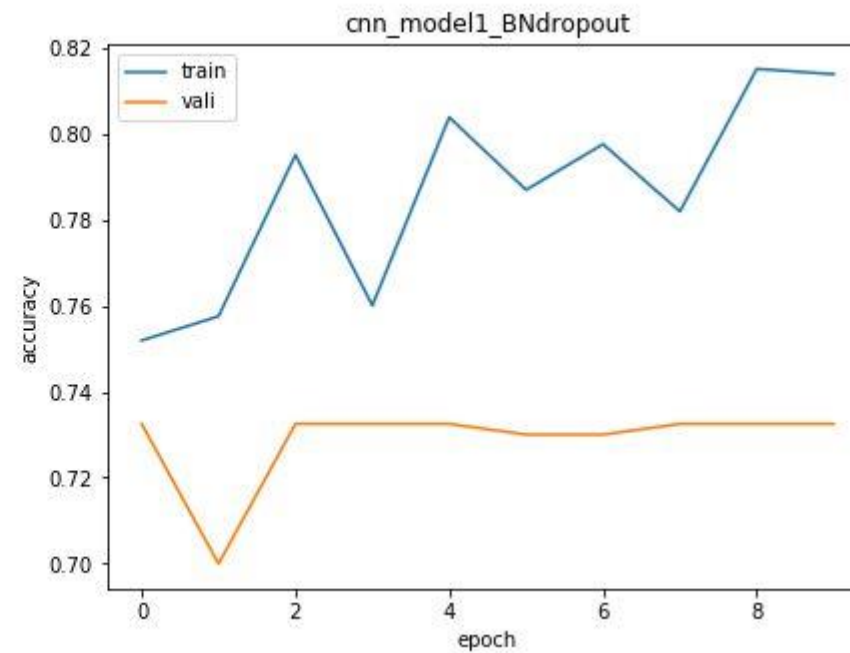
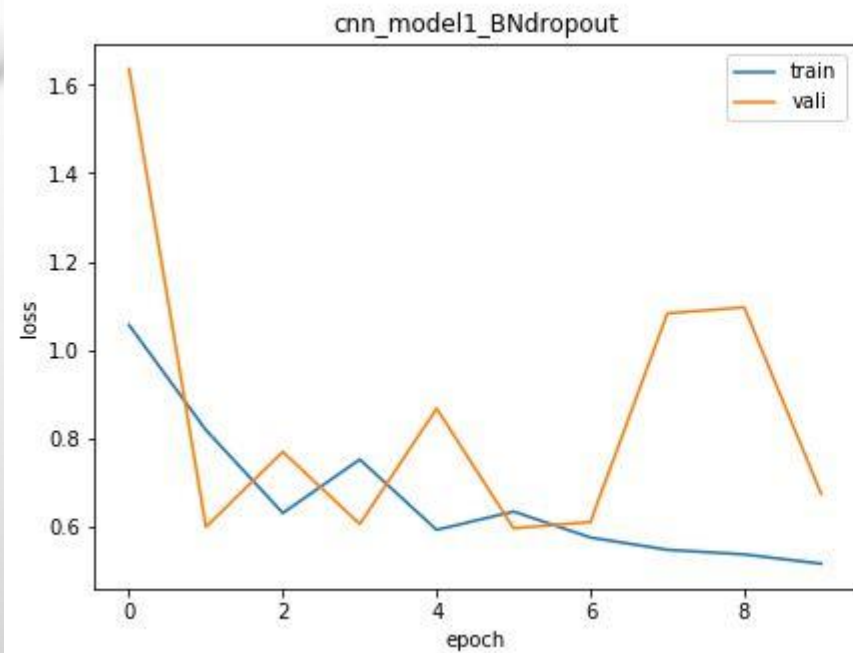
# KERAS NEURAL NETWORK\_1 RELU, 1 SOFTMAX



Without dropout

```
Epoch 1/10  
1550/1600 [=====>.] - ETA: 3s - loss: 2.6622 - acc: 0.8297 Epoch 00000: val_acc improved from -inf to  
0.73250, saving model to model1_dropout_0_00_0.73.hdf5  
1600/1600 [=====] - 126s - loss: 2.6193 - acc: 0.8325 - val_loss: 4.3116 - val_acc: 0.7325
```

# BATCH NORMALIZATION



## Dropout & batch normalization -


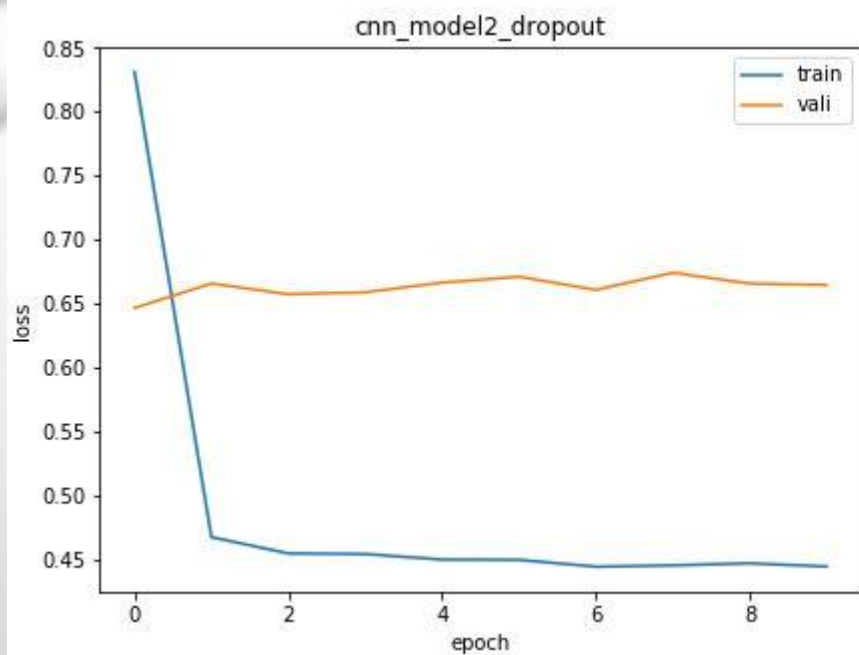
Epoch 9/10

1550/1600 [=====>.] - ETA: 4s - loss: 0.5412 - acc: 0.8123 Epoch 0008: val\_acc did not improve

1600/1600 [=====] - 143s - loss: 0.5356 - acc: 0.8150 - val\_loss: 1.0959 - val\_acc: 0.7325





# KERAS NEURAL NETWORK\_3 CONVO, 1 OUTPUT



## Without dropout

```
Epoch 2/10  
1550/1600 [=====>.] - ETA: 4s - loss: 0.4626 - acc: 0.8342 Epoch 00001: val_acc did not improve  
1600/1600 [=====] - 175s - loss: 0.4675 - acc: 0.8331 - val_loss: 0.6657 - val_acc: 0.7325
```

# BATCH NORMALIZATION



## Dropout and batch normalization

```
Epoch 10/10  
1550/1600 [=====>.] - ETA: 5s - loss: 0.4480 - acc: 0.8323 Epoch 00009: val_acc did not improve  
1600/1600 [=====] - 197s - loss: 0.4467 - acc: 0.8331 - val_loss: 0.6797 - val_acc: 0.7325
```


# STEPS FORWARD

- 1) ADDRESS CLASS IMBALANCE – ACCURACY IS IMPROVED JUST A LITTLE OVER BASELINE
  - 2) DEEPER NETWORKS USUALLY LEARN BETTER – NOT ABLE TO RUN DEEPER NETWORKS WITH CURRENT HARDWARE
- # NEED BIG DATA TOO.....2000 IMAGES ISN'T GOOD ENOUGH DATA FOR NEURAL

# LOGISTIC


```
plt.title('Histogram of Y')  
...: n, bins, patches = plt.hist(Y[:,], 50, facecolor='green', alpha=0.75)
```

Histogram of Y



```
In [99]: plt.title('Histogram of LR_predictions')  
...: n, bins, patches = plt.hist(LR_predictions[:,], 50, facecolor='green', alpha=0.75)
```

Histogram of LR\_predictions



Confusion Matrix: `confusion_matrix(y_true , y_pred)`


```
array([[1624, 2],  
       [ 21, 353]], dtype=int64)
```

```
Accuracy = ((1624+353)/2000)*100  
...: print('Accuracy =' + str(Accuracy) + "%")
```


**Accuracy =98.85 % on Training Data**

# LOGISTIC

```
In [34]: plt.title('Histogram of Y_test')
...: n, bins, patches = plt.hist(Y_test[:,], 50, facecolor='green', alpha=0.75)
```



```
In [35]: plt.title('Histogram of Y_test_predictions')
...: n, bins, patches = plt.hist(Y_test_predictions[:,], 50, facecolor='green', alpha=0.75)
```



Confusion Matrix: `confusion_matrix(y_true , y_pred)`

```
array([[109, 11],
       [ 22,  8]], dtype=int64)
```

```
Accuracy = ((109+8)/150)*100
...: print('Accuracy = ' + str(Accuracy) + "%")
```

**Accuracy = 78.0 % on Test data**

**Upsampled the data**

Original dataset shape Counter({0: 1626, 1: 374})

Resampled dataset shape Counter({0: 1626, 1: 1626})

```
array([[1626,  0],
       [  0, 1626]], dtype=int64)
```

**Accuracy on upsampled data = 100 % on Test data**

```
array([[108, 12],
       [ 23,  7]], dtype=int64)
```

**Accuracy of upsampled model on Test data = 76.66 %**

# QUESTIONS ?